
Quality/Latency-Aware Real-Time Scheduling of Distributed Streaming IoT Applications

Kamyar Mirzazad Barijough, Zhuoran Zhao,
Andreas Gerstlauer

System-Level Architecture and Modeling (SLAM) Lab
Department of Electrical and Computer Engineering
The University of Texas at Austin

<https://slam.ece.utexas.edu>



The University of Texas at Austin

Electrical and Computer Engineering

Cockrell School of Engineering

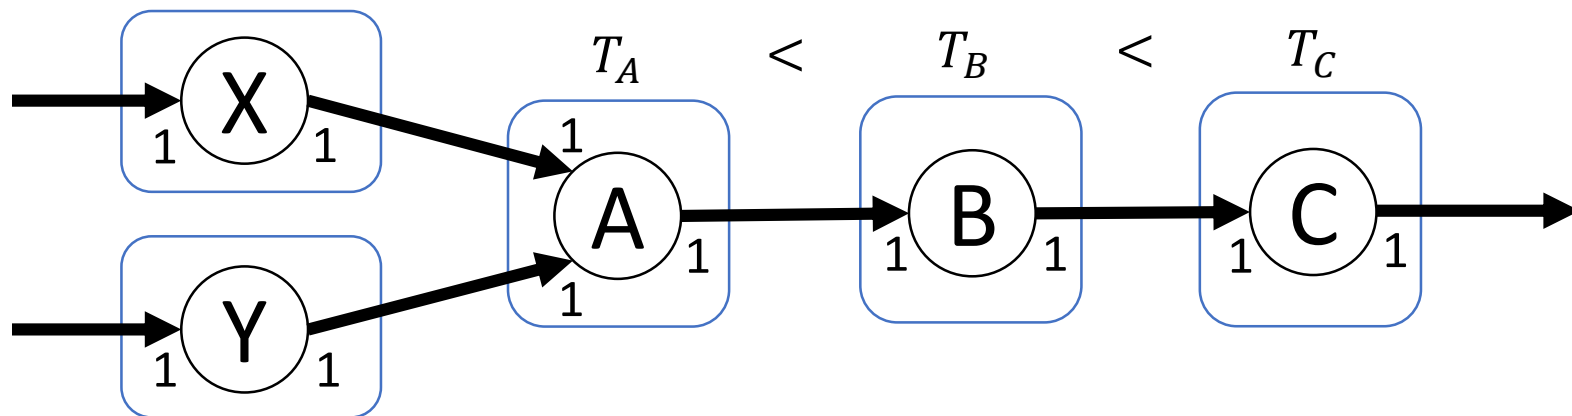


Background

- **Networked and distributed embedded systems**
 - Internet of Things (IoT) and edge computing
 - Networked cyber-physical systems (CPS)
 - Distributed embedded computing
- **Real-time guarantees**
 - Interact with physical world
 - Hard latency requirements
- **Open networks**
 - Dynamically changing traffic sources & patterns
 - Non-deterministic and potentially unbounded latency
- **Key challenge**
 - How to provide real-time guarantees over unpredictable networks?

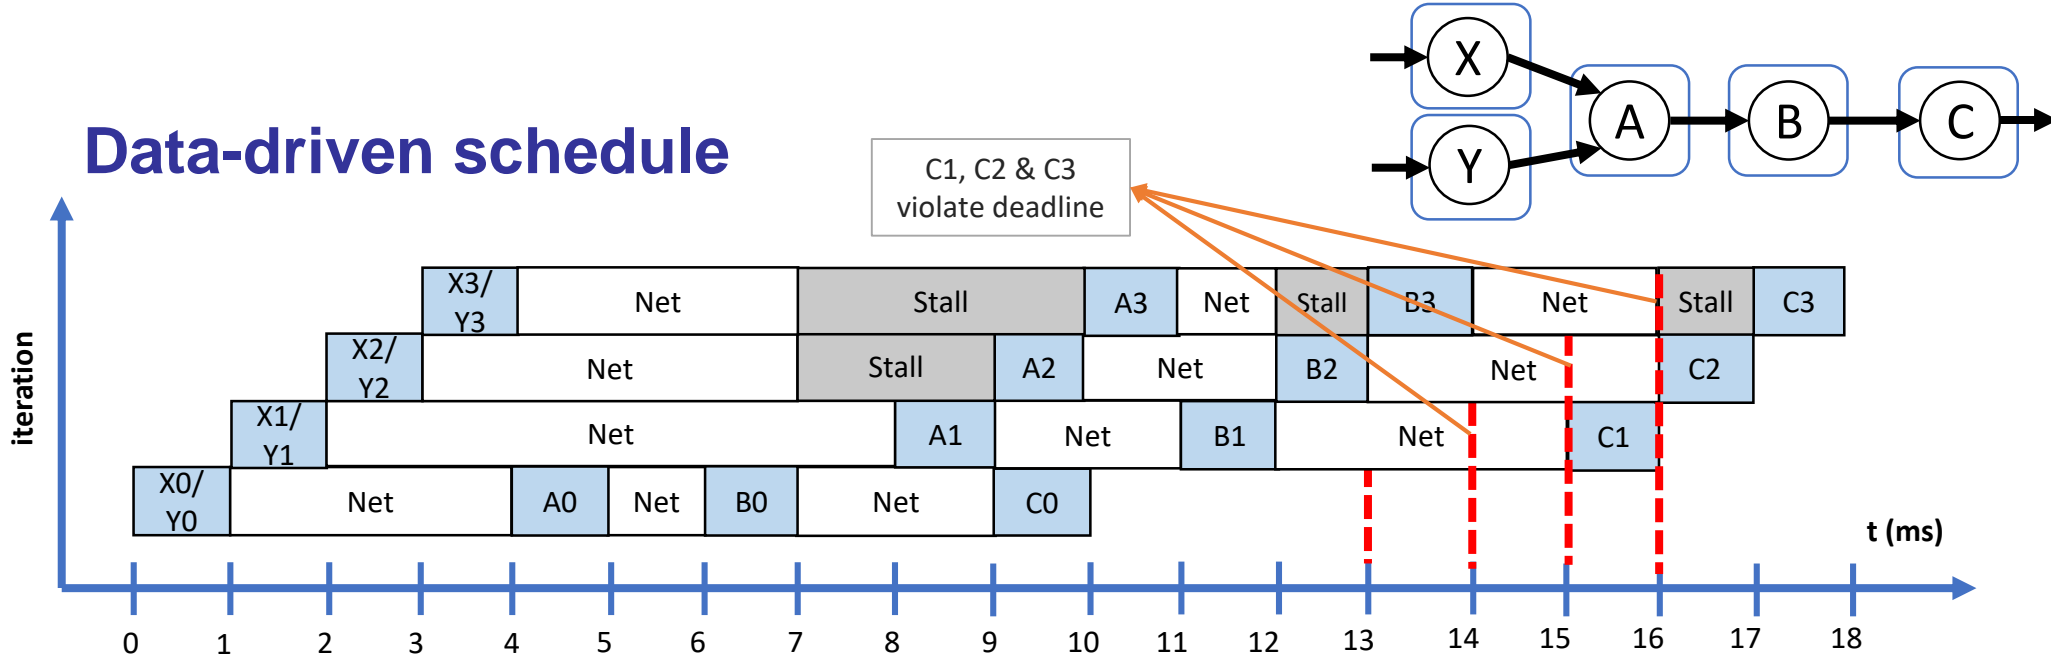
Motivation

- **Embedded applications are often of streaming nature**
 - Best expressed as a data flow graph
- **Latency guarantees provided via timeouts**
 - Tradeoff between latency and losses (quality)
 - Per-actor timeouts
- **Timeout assignment for distributed real-time data flow**
 - Partition latency budget across nodes
 - Application/network-dependent tradeoff




Latency Budget Assignment

Data-driven schedule




Schedule with timeouts & uniform latency distribution

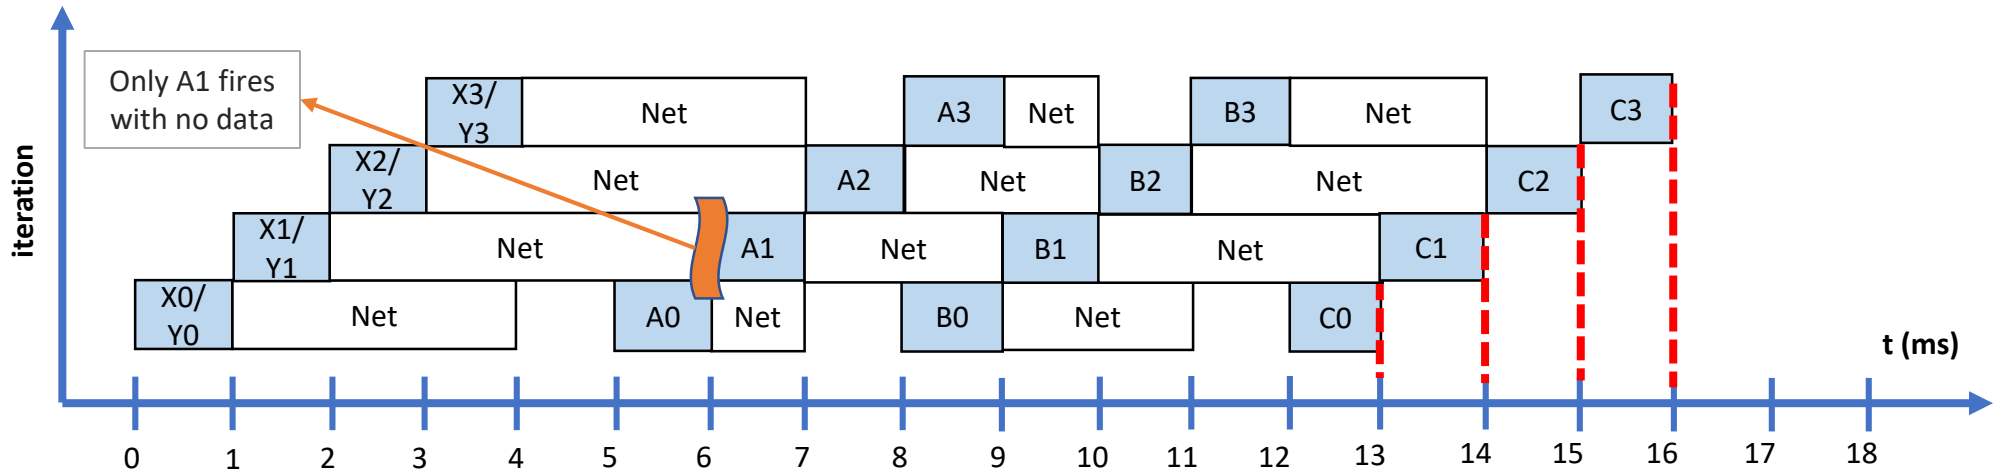


Latency Budget Assignment

Data-driven schedule



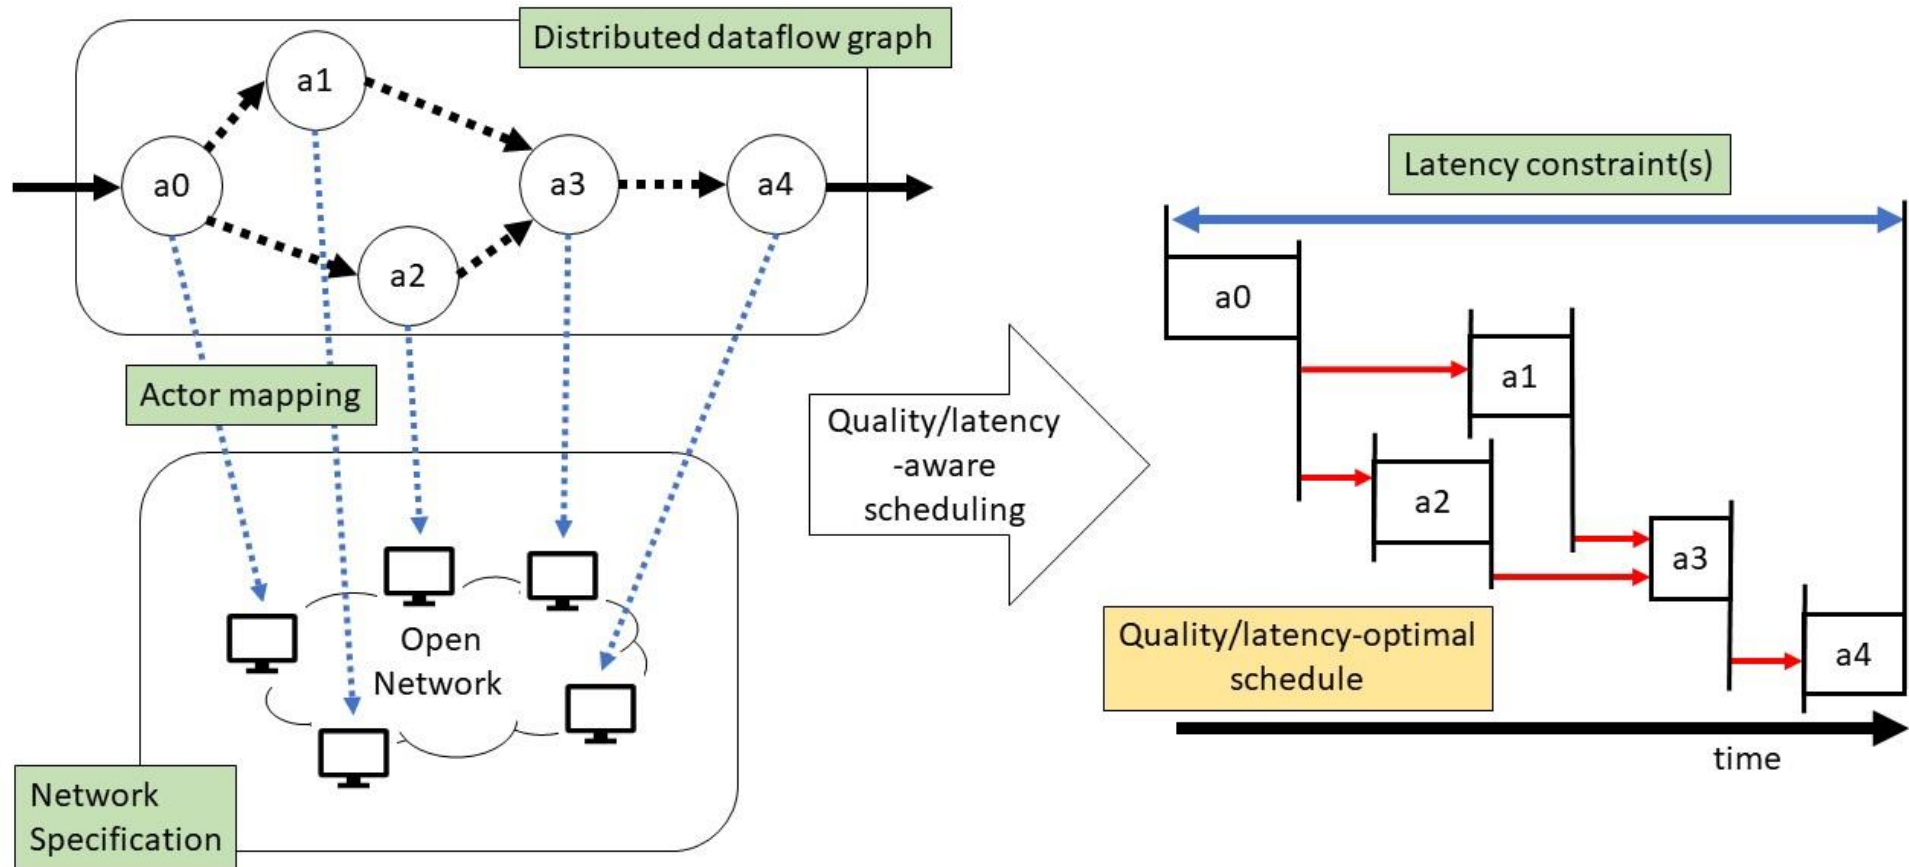
Schedule with optimized latency budget distribution



Related Work

- **Real-time transfer protocol (RTP) [Schulzrinne'03]**
 - Designed for end-to-end data transfer
 - Only pair-wise/end-to-end timeout assignment
- **Distributed real-time computing frameworks**
 - Real-time extension(s) of RPC frameworks [RT-Corba]
 - Stream processing frameworks [Typhoon, Ares, Storm]
 - Requires QoS guarantees or reliable delivery from network
- **PTIDES [Zhao'07]**
 - Discrete-event execution for distributed systems
 - Requires accurate time synchronization and bounded network delay
- **Reactive and Adaptive Data Flow Model (RADF) [Francis'17]**
 - Dataflow with extensions for modeling network effects
 - No timeout assignment/implementation

Overview




- **Derive a schedule for the given data flow graph using**
 - Worst-case execution times (WCET)
 - Mapping information
 - Latency constraints
 - Network specification

Outline

- ✓ **Introduction**
 - ✓ Motivation, background
 - ✓ Related work
- **Formalizing distributed data flow**
 - Timed extension of RADF
- **Scheduling distributed data flow**
 - Quality model and optimization
- **Experimental Results**
- **Summary & Future Work**


Reactive and Adaptive Data Flow (RADF)

- **Model data losses in network channels as “empty” tokens**
 - Maintain deterministic execution in presence of losses
- **Channels could be lossless or lossy**
 - Traditional channels are “lossless”
 - “Lossy” channels can make a token “empty”
- **Actors need to handle network losses**
 - Can consume empty token(s), but produce only non-empty tokens
 - Multiple firing rules based on input token patterns



Timed Extension of RADF (T-RADF)

- **Empty tokens need to be injected by the runtime**
 - Decide based only on local time
 - Relative timeouts between firings
- **T-RADF extends RADF with rates on input and outputs**
 - Set (average) timeouts based on firing rates
 - Firing rates derived from external rates + repetition vector



Schedule Computation

- **Assumptions**

- Homogeneous T-RADF
 - Any graph can be made homogeneous albeit exponentially larger
- One actor per host
 - Statically schedule actors mapped to the same host into a super-actor
 - Might lead to deadlock, CSDF can relax this (future work)

- **Conservative analysis**

- Fixed static schedule with specified periods
- Can adjust schedule dynamically to optimize latency/quality
 - Derive relative start time offsets/phase shifts

Schedule Computation

- **Latency l between input-output pair**

- Depends on actor execution times e_i and channel delays d_j


$$l = \sum_{i=0}^{m-1} e_i + \sum_{j=1}^{m-1} d_j \leq l'$$

- **Latency constraint l' requires bounding e_i and d_j**

- Worst-case execution time bounds: $e_i \leq e'_i$
- Goal: find bounds d'_j for d_j

- **Find d'_j to satisfy l' and maximize output quality Q**

- d'_j affects token delivery probability p_j and therefore quality
- Quality model to describe Q in terms of d'_j



Quality Model (1)

- **SNR as quality metrics**

- Signal processing applications
- Quantify noise power of the output

- **Single lossy channel**


- Express token delivery probability p_g as function of delay budget d'_g

$$p_g(d'_g) = (1 - \mu_g) \cdot F_{D_g}(d'_g)$$

with network parameters:

μ_g : average loss rate of channel c_g

$F_{D_g}(d_g)$: cumulative distribution function (CDF) of random network delay D_g



Quality Model (2)

- **Single lossy channel (cont'd)**

- Noise in channel c_g at iteration i : $n_g[i] = |s_g[i] - R_g(s_g)|$

$s_g[i]$: (original) signal value of c_g at iteration i

R_g : replacement function of actor a_g

- Assuming n_g is upper bounded by s_g : $n_g[i] \leq \alpha s_g[i]$


- **Actor chain**

- Assuming system and, hence, actors are linear with weights w_j

$$n_m[i] \leq \alpha \tilde{w} s_0[i] \quad \text{where } \tilde{w} = \prod_{j=1}^m w_j$$

- This equation holds regardless of number of losses

➤ **Calculate minimum SNR from maximum noise power**



Quality Model Computation

- **In graph with more than one actor chain**
 - Iterate over paths and accumulate noise power
 - Multiple iterations for cyclic graphs
 - Exponential complexity
- **A better alternative**
 - Breadth-first traversal of channels
 - Accumulate partial noise power at each channel
 - $O(I * m^2 / n)$ complexity
 - in number of inputs I , channels m and nodes n

Scheduling Formulation

- **Optimization problem: find the schedule that maximizes quality**

$$\underset{\mathbf{d}'}{\text{maximize}} \quad Q(\mathbf{d}', \tau)$$

$$\text{subject to} \quad d'_j \geq 0, \quad \forall j \in \text{channels}. \quad \text{precedence constraints}$$

$$\sum_{i=0}^{m_k-1} e'_{(k,i)} + \sum_{j=1}^{m_k-1} d'_{(k,i)} \leq l'_k, \quad \forall k \in \text{paths}. \quad \text{latency constraints}$$

- **Solving the optimization problem**
 - Q is non-linear/non-convex but has closed form and is differentiable
 - d' are continuous
 - Use numerical gradient-based iterative methods
 - E.g. Constrained Trust Region (CTR) solver

Outline

- ✓ **Introduction**
 - ✓ Motivation, background
 - ✓ Related work
- ✓ **Formalizing distributed data flow**
 - ✓ Timed extension of RADF
- ✓ **Scheduling distributed data flow**
 - ✓ Quality model and optimization
- **Experimental Results**
- **Summary & Future Work**

Experimental Setup (1)

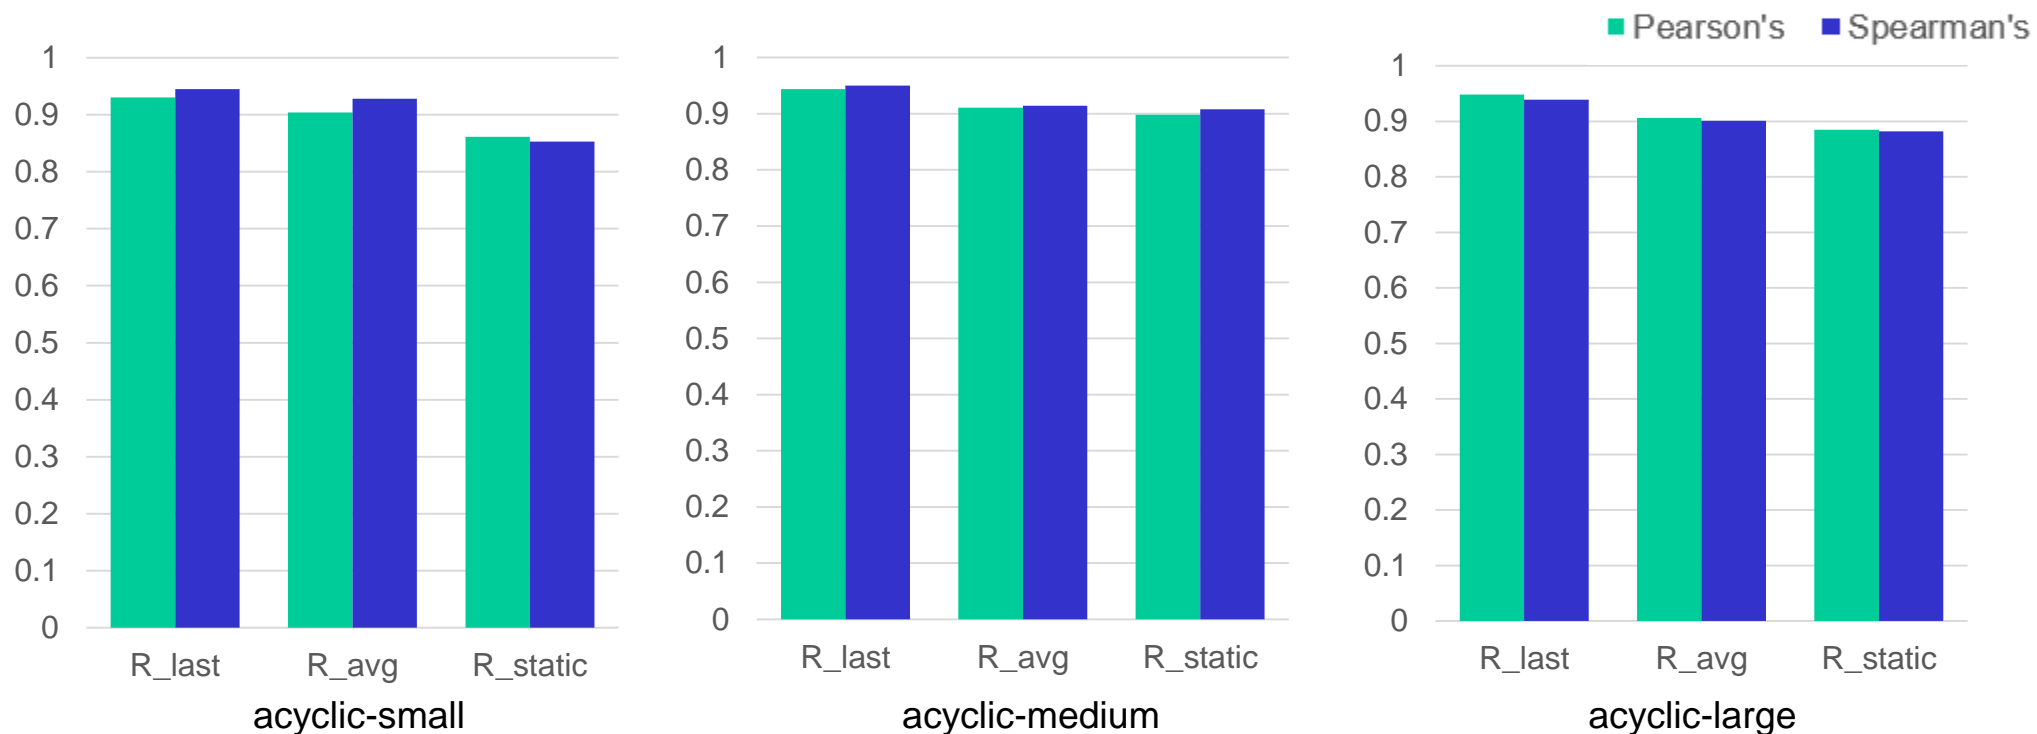
- **T-RADF and scheduler implemented in Python**
 - NetworkX for graph representation
 - Scipy for optimization with 10^{-5} as gradient norm threshold
 - 5 iterations for cyclic graphs
- **Application models**
 - 100 random cyclic/acyclic graphs with 10/50/100 nodes [sdf3]
 - Three replacement policies
 - R_{static} : replaces empty tokens with zeros
 - R_{last} : replaces empty tokens with the last received value
 - R_{avg} : replaces with the running average of received values

Experimental Setup (2)

- **Simulated via OMNET++ and INET Framework**
 - UDP sockets for lossy channels
 - INET's cloud model (5Mbps, $\mu = 1\%$, Gamma NDD)
- **Relative latency constraints**
 - l_{min}/l_{max} are latencies for delivery probabilities of 0.1% and 99.9%
 - Constraint factor $\rho = (l' - l_{min}) / (l_{max} - l_{min})$

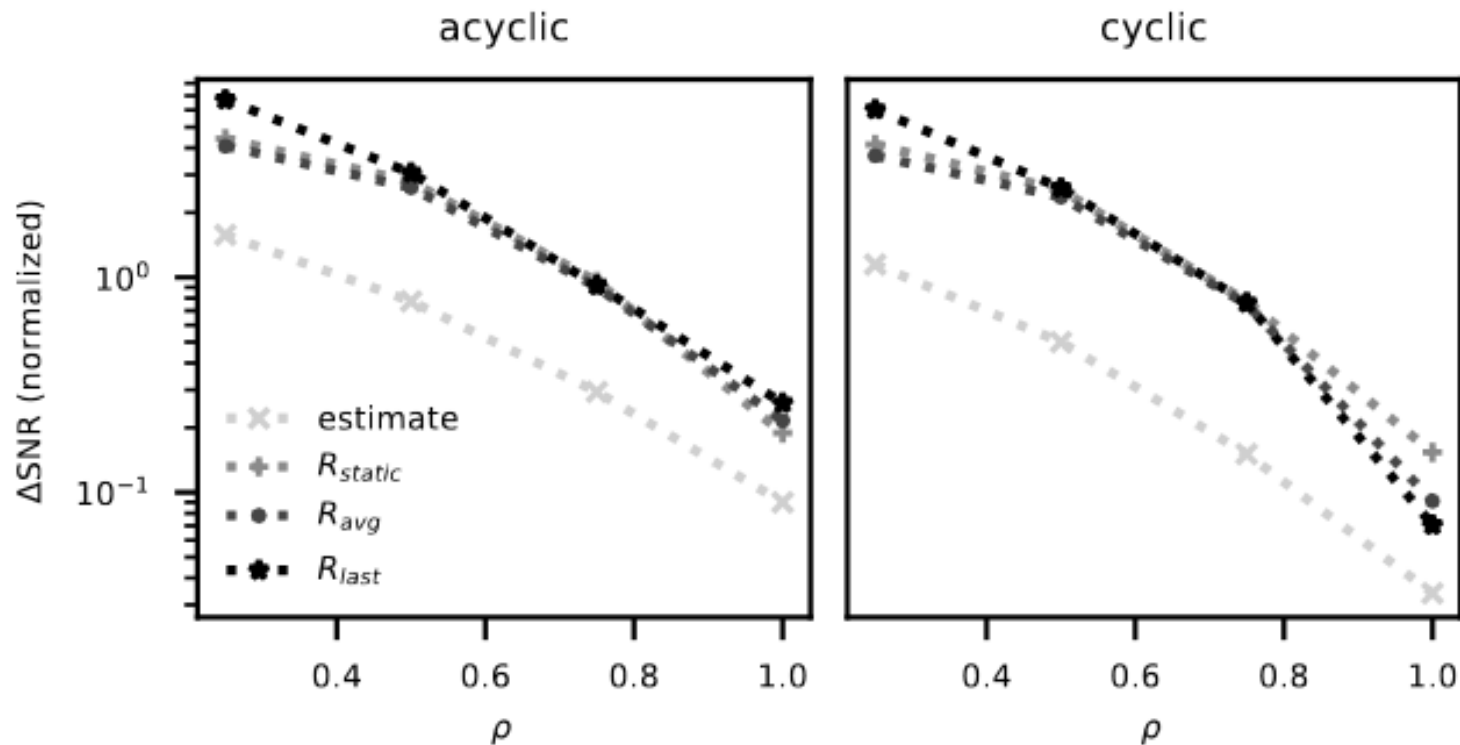
Quality Model Fidelity

- **Correlation between estimated and measured SNR**
 - Relatively/monotonicity: Spearman's correlation coefficient
 - Absolute values: Pearson's correlation coefficient
- **Setup**
 - Generated 100 random schedules for ten graphs
 - Latency constraint factor ρ randomly chosen in the interval $[0.1, 0.9]$



Optimization Results

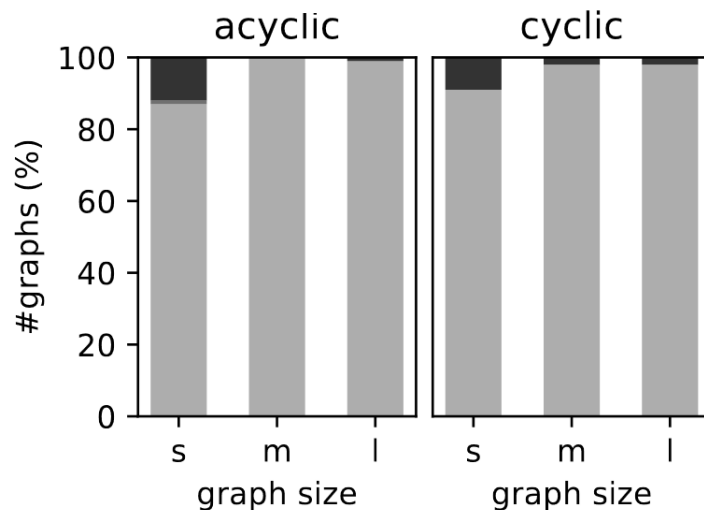
- **Measured and predicted SNR improvement**
 - Vs. uniform budget distribution as baseline schedule
 - Averaged across sizes



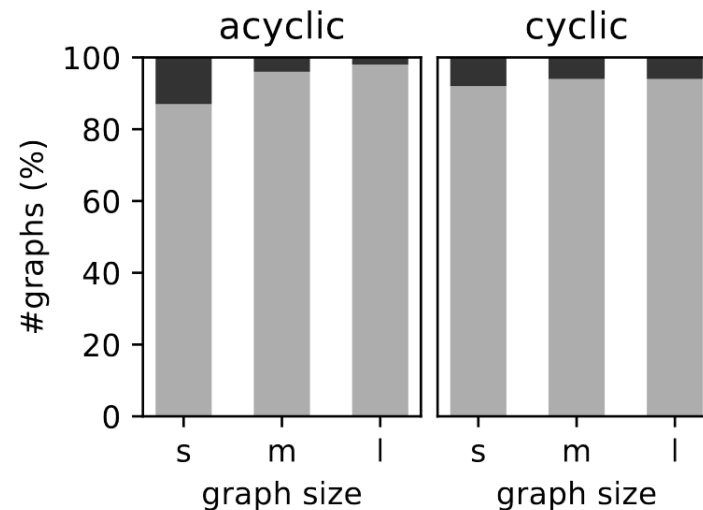
- **Higher gains with tight constraints**
 - Up to 900%, on average 75%
- **Conservative estimate from quality model**

Optimization Outcomes


- Optimization success rate




(a) $\rho=0.25$




(b) $\rho=0.5$



(c) $\rho=0.75$

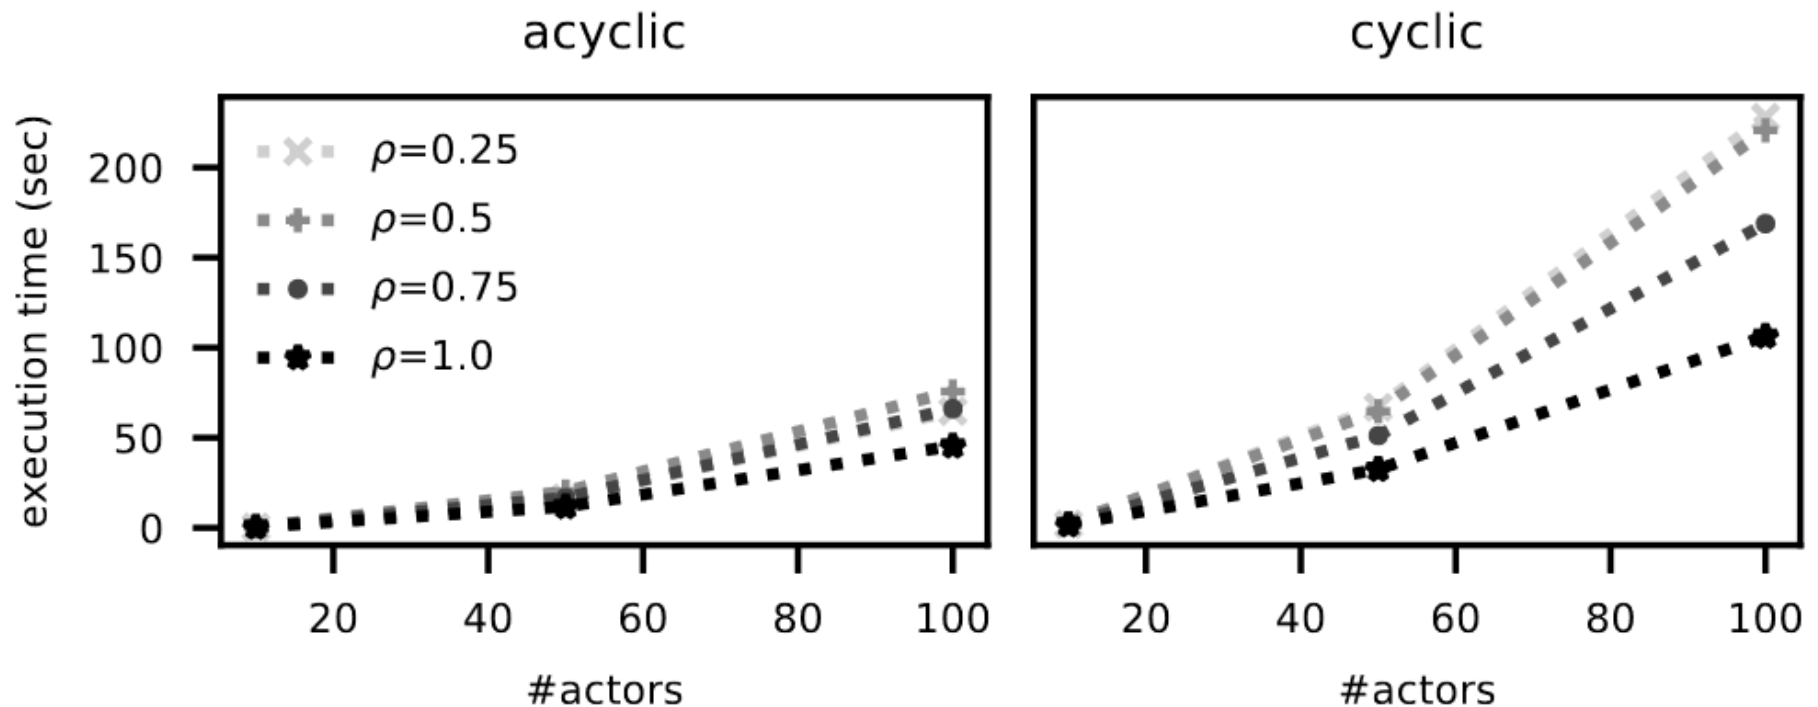


(d) $\rho=1.0$



Optimization Runtime


- **Average runtime of optimization solver**
 - Measured on Intel Core i7-920



- **Larger for tight constraints**
- **Larger for cyclic graphs due to multiple iterations**

Distributed Neural Network Example

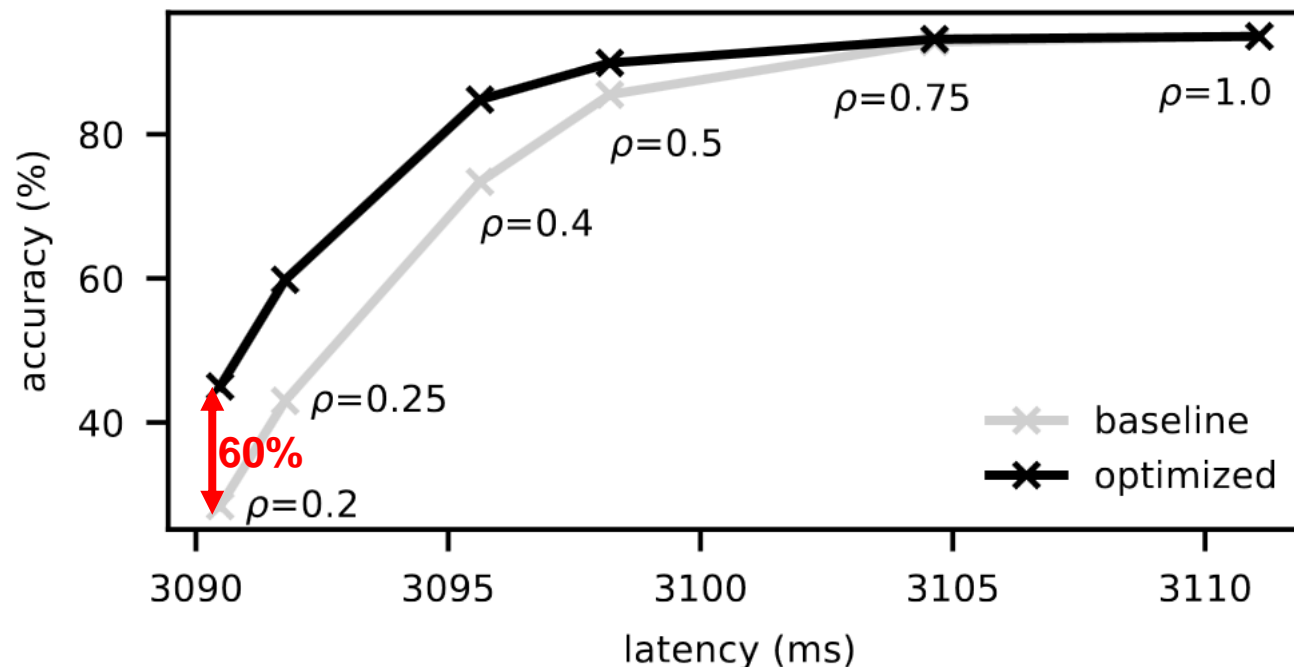
- **Two-layer network for MNIST digit recognition**
 - Non-linear due to activation functions
 - Each token is 16 doubles
 - WCET of 1sec for FC1/FC2
 - Account for network delay of 49 tokens in WCET of Input



- **Significant accuracy gains under tight latency constraints**
 - Up to 60%, on average 20%

Distributed Neural Network Example

- **Two-layer network for MNIST digit recognition**
 - Non-linear due to activation functions
 - Each token is 16 doubles
 - WCET of 1sec for FC1/FC2
 - Account for network delay of 49 tokens in WCET of Input



- **Significant accuracy gains under tight latency constraints**
 - Up to 60%, on average 20%

Summary and Future Work

- **Quality/latency-aware scheduling**
 - Dataflow models for distributed embedded systems
 - Quality/latency tradeoff
 - Scheduling dist. dataflow by optimizing the tradeoff
 - Tools, graphs and simulation models available online
<https://github.com/SLAM-Lab/QLA-RTS>
- **Future work**
 - Address the restrictions
 - Homogeneity, one-actor-per-host mapping
 - Runtime system for T-RADF
 - Dynamic scheduling